

Report of Operating System Curriculum Design

叶剑飞

Victor Jianfei Ye

September 7, 2012

Contents

1	Expense Comparison among different I/O Functions	2
1.1	File I/O Functions Supported by ANSI C	2
1.2	File I/O Functions Supported by UNIX C	2
1.3	Memory Map	2
2	Make a Simple Shell	4
2.1	Victor Shell	4
3	Synchronization among Processes/Threads	5
3.1	Synchronization among Processes	5
3.2	Synchronization among Threads	6
A	Source Code of File I/O Operation	7
A.1	file1.c	7
A.2	file2.c	10
A.3	file3.c	12
B	Source Code of a Simple Shell	14
B.1	victorshell.c	14
C	Source Code of Synchronization	24
C.1	process_synchronization.c	24
C.2	thread_synchronization.c	29

1 Expense Comparison among different I/O Functions

There are three methods of file I/O operation in Linux operating system, namely ANSI C I/O functions, UNIX C I/O functions and memory map. We now make a file reversion operation to test the time consumed by the three kinds of functions. All the memories are allocated dynamically. By the way, my C source code includes “for” loop initial declarations. This is not allowed according to C89 standard, but C99 standard supports it. Consequently, “-std=c99” is needed while compiling my C source code if you are using “gcc”.

1.1 File I/O Functions Supported by ANSI C

The file I/O functions supported by ANSI C are `fopen`, `fclose`, `fread`, `fwrite`, `fprintf`, `fscanf`, `fgets`, `fputs`, etc. What we need now is binary operation. Consequently, only `fopen`, `fclose`, `fread`, `fwrite` are used.

Firstly, `fopen` is used to open a file. On success, it returns a pointer to the file structure (`FILE`). Otherwise, `NULL` is returned. After that, we can read the test file by the `fread` function. We read the file n characters¹ every time. Then the array will be reversed. The reversed array should be written back into the original file by the `fwrite` function. Writing speed is still n characters every time. Finally the file should be closed by the `fclose` function.

1.2 File I/O Functions Supported by UNIX C

The file I/O functions supported by UNIX C are very similar in forms to those supported by ANSI C. However, there are only binary operations. Firstly, the file should be open by the `open` function. On success, it returns an integer, which is the file descriptor. Otherwise, `-1` is returned. Then the `read` function can be used to read the file data. The characters read each time is still from the second command-line argument vector. The array should also be reversed and write back to the original file by the `write` function. Finally, the file descriptor should be closed by the `close` function.

1.3 Memory Map

The `mmap` function map files into memory. While we are operating the memory, we are actually modifying the file. We still need to open the file by the `open` function. The `mmap` function needs the file descriptor. It is used as below:

```
mmap( NULL, file size , PROT_READ|PROT_WRITE,  
      MAP.SHARED, fd , 0 );
```

After mapping, only thing we need is to reverse the array in the memory. But we are actually reversing the file. Finally unmap the file and close the file descriptor.

¹ n comes from the second command-line argument vector, i.e. `argv[2]`.

According to the test, memory map is the most effective way to read and write the file.

2 Make a Simple Shell

2.1 Victor Shell

My English name is Victor. As a result, the shell made by me is called “Victor Shell” and an alternative name “victorshell”, with the copyright information in the “version” part, which is “Copyright © Victor Jianfei Ye”².

A “shell” is nothing but a command interpreter. After pressing the “Enter” key, the command should be execute. The shell need to analyze the command string. I analyze the command string by the `sscanf` function in order to seperate the space and tabulate key. Then the shell should judge whether it is a internal command. If it is, then the shell should execute the command by the shell itself. Otherwise, the command should be regarded as external command and create a new process to execute the command. If it failed, it means that it isn’t a external command and therefore the shell should prompt the error message “command not found”. The command could be read from either keyboard (stdin) or a shell script file. It depends on the command-line argument to the shell. As a result, the command is read by `fgets` function. The internal command supported by my shell are “exit”³, “cd”⁴, “echo”⁵, “pwd”⁶, “clear”⁷, “environ” and “help”. Another thing I did is to launch a background application. When “&” is typed after the command. The program should be launched immediately but run in background. As a result, I need to judge whether the last character is “&”. If it is, I just simply assign “false” the boolean variable *background*. The shell launch the application as usual, but the shell need to judge whether it need to block itself to wait for the child process according to the boolean variable *background*. If the value is “true”, then wait for the child process by the function `waitpid`. Otherwise, the shell cannot wait for the child process in the current thread, but it need to create a new thread in order to wait for the child process in order to prevent the defunct process. Consequently, my C source code should be linked to the POSIX thread library `pthread` while linking the object files, by typing the command-line linking argument “-lpthread”. By the way, compiling my C source code of “victorshell” also needs the command-line argument “-std=gnu99”.

²“Jianfei Ye” is my Chinese name (叶剑飞) in Pinyin, written in the English style, i.e. the given name is before the family name

³Exit the shell with the return value, default to zero (i.e. `EXIT_SUCCESS`).

⁴Change the current working directory. Support “cd - ” and “cd ~”, even bare “cd”.

⁵Simply print the words after “echo”, without “\$” variable support.

⁶Simply print the current working directory.

⁷Simply clear the screen by a series of control characters, that is “\033[H\033[2J”.

3 Synchronization among Processes/Threads

3.1 Synchronization among Processes

Synchronization among processes in Linux operating system needs the System V semaphore. It contains three important APIs, namely, `semget`, `semctl` and `semop`. These three APIs cost me many days to comprehend.

The `semget` function is used to create a semaphore. The first parameter is the key. We can simply use `IPC_PRIVATE` here. The second parameter is the amount of the semaphores in the semaphore set. The third parameter is the flag to the semaphore, specifies both `IPC_CREAT` and `IPC_EXCL` and a semaphore set already exists for key. On success, it returns the semaphore set identifier, otherwise -1 is returned.

The `semctl` function performs the control operation to the semaphore. The first parameter is the semaphore set identifier, which is from the return value of the function `semget`. The second parameter is the subscript to the semaphore set. The third parameter is the exact operation command, e.g. `SETVAL`, `IPC_RMID`, etc. The fourth parameter is a union type variable. This function returns -1 on failure.

The `semop` function performs operations on selected semaphores in the set. The first parameter is the semaphore set identifier. The second parameter is a pointer to the structure type “struct sembuf”, which contains the semaphore number, semaphore operation and operation flags. The third parameter refers to the amount of elements in the array.

Now we know the meaning of the three System V APIs. We can write P/V operations by these functions, as the code shown below:

```
bool P( int semid )
{
    struct sembuf spos;
    spos.sem_num = 0;
    spos.sem_op = -1;
    spos.sem_flg = 0;
    if ( semop( semid, &spos, 1 ) == -1 )
        return false;
    return true;
}
```

```
bool V( int semid )
{
    struct sembuf spos;
    spos.sem_num = 0;
    spos.sem_op = 1;
    spos.sem_flg = 0;
    if ( semop( semid, &spos, 1 ) == -1 )
        return false;
}
```

```
        return true ;  
    }
```

The return value is boolean type, which stands for whether the operation succeed or not.

3.2 Synchronization among Threads

Thread synchronization is much easier than process synchronization. Only sem_wait function and sem_post function are needed. Before that, we need the function sem_init to initialize the semaphore. Be careful, it needs POSIX thread library, so link with the command-line argument “-lpthread”.

All of my C source codes in process synchronization and thread synchronization have “for” loop initial declarations, which is only supported by C99 standard. Moreover, I also used “usleep” function in order to suspend execution for microsecond intervals, but it isn’t an ANSI C function. Thus, the source codes should be compiled by the command-line argument “-std=gnu99”.

A Source Code of File I/O Operation

A.1 file1.c

```
/*
 * file1.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *                 gcc file1.c -o file1 -std=c99 -pedantic -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>

int main ( int argc , char * argv[] )
{
    FILE * fp = NULL;
    char * pText = NULL;
    char temp;
    int bytes , offset , filesize;
    struct stat statbuf;
    if ( argc != 3 )
    {
        fprintf( stderr , "Usage: %s <filename> <bytes>\n\n" , argv[0] );
        return EXIT_FAILURE;
    }
    if ( sscanf( argv[2] , "%d" , &bytes ) != 1 )
    {
        fprintf( stderr , "Bytes number must be number!\n\n" );
        return EXIT_FAILURE;
    }
    if ( stat( argv[1] , &statbuf ) != 0 )
    {
        fprintf( stderr , "File Access Error\n\n" );
        return EXIT_FAILURE;
    }
    filesize = statbuf.st_size;
```

```

fp = fopen( argv[1], "rb" );
if ( fp == NULL )
{
    fprintf( stderr, "File_Open_Error\n\n" );
    return EXIT_FAILURE;
}
pText = (char *)malloc( filesize * sizeof(char) );
if ( pText == NULL )
{
    fclose(fp);
    fprintf( stderr, "No_Enough_Memory!!!\n\n" );
    return EXIT_FAILURE;
}
offset = 0;
while ( !feof(fp) )
{
    fread( pText+offset, bytes * sizeof(char), 1, fp );
    offset += bytes;
}
fclose(fp);
fp = NULL;
for ( int i = (filesize >> 1) ; i < filesize; i ++ )
{
    temp = pText[i];
    pText[i] = pText[filesize-i-1];
    pText[filesize-i-1] = temp;
}
fp = fopen( argv[1], "wb" );
if ( fp == NULL )
{
    fprintf( stderr, "File_Open_Error\n\n" );
    free( pText );
    return EXIT_FAILURE;
}
offset = 0;
while ( offset+bytes < filesize )
{
    fwrite( pText+offset, bytes * sizeof(char), 1, fp );
    offset += bytes;
}
fwrite( pText+offset, (filesize-offset)*sizeof(char), 1, fp );
fclose(fp);
free(pText);
return EXIT_FAILURE;

```


}

A.2 file2.c

```
/*
 * file2.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *                 gcc file2.c -o file2 -std=c99 -pedantic -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>

int main ( int argc , char * argv[] )
{
    int fd = 0;
    char * pText = NULL;
    char temp;
    int bytes , offset , filesize;
    struct stat statbuf;
    if ( argc != 3 )
    {
        fprintf( stderr , "Usage: %s <filename> <bytes>\n\n" , argv[0] );
        return EXIT_FAILURE;
    }
    if ( sscanf( argv[2] , "%d" , &bytes ) != 1 )
    {
        fprintf( stderr , "Bytes number must be number!\n\n" );
        return EXIT_FAILURE;
    }
    fd = open( argv[1] , ORDWR );
    if ( fd == -1 )
    {
        fprintf( stderr , "File Open Error\n\n" );
        return EXIT_FAILURE;
    }
    if ( fstat( fd , &statbuf ) != 0 )
    {
```

```

        fprintf( stderr , "Error_Getting_File_Information\n\n" );
        close(fd);
        return EXIT_FAILURE;
    }
    filesize = statbuf.st_size;
    pText = (char *)malloc( filesize * sizeof(char) );
    if ( pText == NULL )
    {
        close(fd);
        fprintf( stderr , "No_Enough_Memory!!!\n\n" );
        return EXIT_FAILURE;
    }
    offset = 0;
    while ( offset < filesize )
        offset += read( fd, pText+offset, bytes * sizeof(char) );
    if ( lseek( fd, 0, SEEK_SET ) == -1 )
    {
        fprintf( stderr , "lseek_error\n\n" );
        free( pText );
        close(fd);
        return EXIT_FAILURE;
    }
    for ( int i = (filesize >> 1) ; i < filesize; i ++ )
    {
        temp = pText[i];
        pText[i] = pText[filesize-i-1];
        pText[filesize-i-1] = temp;
    }
    offset = 0;
    while ( offset < filesize )
        offset += write( fd, pText+offset, bytes*sizeof(char) );
    close(fd);
    free(pText);
    return EXIT_FAILURE;
}

```

A.3 file3.c

```
/*
 * file3.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *                 gcc file3.c -o file3 -std=c99 -pedantic -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <sys/stat.h>
#include <fcntl.h>
#include <unistd.h>
#include <sys/mman.h>

int main ( int argc , char * argv[] )
{
    int fd = 0;
    char * pText = NULL;
    char temp;
    int filesize;
    struct stat statbuf;
    if ( argc != 2 )
    {
        fprintf( stderr , "Usage: %s <filename>\n\n" , argv[0] );
        return EXIT_FAILURE;
    }
    fd = open( argv[1], ORDWR );
    if ( fd == -1 )
    {
        fprintf( stderr , "File_Open_Error\n\n" );
        return EXIT_FAILURE;
    }
    if ( fstat(fd, &statbuf) != 0 )
    {
        fprintf( stderr , "Error_Getting_File_Information\n\n" );
        close(fd);
        return EXIT_FAILURE;
    }
}
```

```

    filesize = statbuf.st_size;
    pText = mmap( NULL, filesize ,
                  PROT_READ|PROT_WRITE, MAP_SHARED, fd, 0 );
    for ( int i = (filesize >> 1) ; i < filesize; i ++ )
    {
        temp = pText[i];
        pText[i] = pText[filesize-i-1];
        pText[filesize-i-1] = temp;
    }
    munmap( pText, filesize );
    close(fd);
    return EXIT_FAILURE;
}

```

B Source Code of a Simple Shell

B.1 victorshell.c

```
/*
 * victorshell.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *     gcc victorshell.c -o victorshell \
 *     -std=gnu99 -D_REENTRANT -lpthread -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <string.h>
#include <stdbool.h>
#include <stdarg.h>
#include <sys/types.h>
#include <pwd.h>
#include <sys/wait.h>
#include <pthread.h>

#define MAXLENGTH 256

typedef int COUNT;

typedef struct
{
    pid_t pid;
    char cmdline[MAXLENGTH];
} PROCESS_STRUCTURE;

void version(void)
{
    puts( "Victor Shell , version 1.0.0 - beta" );
    puts( "Copyright (C) 2012 Victor Jianfei Ye" );
    puts( "" );
}
```

```

void arghelp( const char * currentFilename )
{
    version();
    printf ( "Usage:%s_[option]_...\n", currentFilename );
    puts ( "Options:" );
    puts( "\t—version");
    puts( "\t—help");
    puts( "\t-v" );
}

void help(const char * content)
{
    if ( content == NULL )
    {
        version();
        puts ( "exit" );
        puts ( "cd" );
        puts ( "clear" );
        puts ( "pwd" );
        puts ( "echo" );
        puts ( "" );
    }
    else
    {
        if ( !strcmp( content , "exit" ) )
        {
            puts("exit:_exit_[n]");
            puts("\tExit_the_shell.");
            puts("");
            puts("\tExits_the_shell_with_a_status_of_N.");
            puts("If_N_is_omitted,_the_exit_status");
            puts("\tis_that_of_the_last_command_executed.");
        }
        else if ( !strcmp( content , "cd" ) )
        {
            puts( "cd:_cd_<directory_name>" );
            puts( "\tchange_to_the_directory" );
            puts( "" );
        }
        else if ( !strcmp( content , "clear" ) )
        {
            puts( "clear:_clear_" );
            puts( "\tclear_the_terminal_screen" );
        }
    }
}

```



```

        exit( exitcode );
    else
    {
        printf("victorshell:_\n"
        printf("exit:_%s:_numeric_argument\n",
            argv[1] );
        exit(255);
    }
}
else
{
    puts("victorshell:_exit:_too_many_arguments\n");
    exit(EXIT_FAILURE);
}
}
else if ( !strcmp(argv[0], "cd" ) )
{
    if( argc == 1 || (argc == 2 && !strcmp( argv[1],
    {
        if ( chdir(getenv("HOME")) != 0 )
            printf( "victorshell:_cd:_no_such_directory\n" );
    }
    else if ( argc == 2 && !strcmp( argv[1] , "-" ) )
    {
        char * oldPwd = getenv("OLDPWD");
        char curPwd[MAXLENGTH];
        if ( getcwd( curPwd, sizeof(curPwd) ) ==
            strncpy( curPwd, "", 3 ) ;
        if ( oldPwd == NULL )
            puts( "victorshell:_cd:_OLDPWD_not_set\n" );
        else
        {
            if ( chdir(oldPwd) != 0 )
                printf( "victorshell:_cd:_no_such_directory\n" );
            else
            {
                puts( oldPwd );
                if ( strcmp( curPwd, "" )
                    setenv( "OLDPWD",
                        else
                            unsetenv( "OLDPWD" );
            }
        }
    }
}
}

```

```

else
{
    char curPwd[MAXLENGTH];
    if ( getcwd( curPwd, sizeof(curPwd) ) ==
        strncpy( curPwd, "", 3 ) ;
    if( chdir(argv[1]) != 0 )
        printf( "victorshell:_cd:_%s:_no_"
    else
    {
        if ( strcmp( curPwd, "" ) )
            setenv( "OLDPWD", curPwd,
        else
            unsetenv( "OLDPWD" );
    }
}
return true;
}
else if ( !strcmp(argv[0], "echo" ) )
{
    if ( argc > 1 )
    {
        printf( "%s", argv[1] );
        for ( int i = 2; i < argc; i ++ )
            printf ( "_%s", argv[i] );
    }
    puts("");
    return true;
}
else if ( !strcmp( argv[0], "pwd" ) )
{
    char path[MAXLENGTH];
    if ( getcwd(path, sizeof(path)) == NULL )
        puts("invalid _path");
    else
        puts(path);
    return true;
}
else if ( !strcmp( argv[0], "clear" ) )
{
    printf("\033[H\033[2J");
    return true;
}
else if ( !strcmp( argv[0], "environ" ) )
{

```

```

        char path[MAX_LENGTH];
        getcwd( path , sizeof(path) );
        printf( "USER=%s\n", getpwuid( getuid() )->pw_name );
        printf( "PWD=%s\n", path );
        printf( "HOME=%s\n", getenv( "HOME" ) );
        printf( "PATH=%s\n", getenv( "PATH" ) );
        return true;
    }
    else if ( !strcmp( argv[0], "help" ) )
    {
        if ( argc == 1 )
            help(NULL);
        else
            help( argv[1] );
        return true;
    }
}
return false;
}

void * waitBackgroundProcess ( void * args )
{
    int status;
    PROCESS_STRUCTURE processStructure = *(PROCESS_STRUCTURE *)args;
    pid_t pid = processStructure.pid;
    waitpid( pid , &status , 0 );
    if ( WEXITSTATUS(status) == EXIT_FAILURE )
    {
        fprintf( stderr , "victorshell: %s: command not found...\n" ,
        )
    }
    else
    {
        printf ( "Done\t\t\t%s\n", processStructure.cmdline );
    }
    return NULL;
}

bool exeExternalCmd( const char * cmd )
{
    int argc;
    COUNT i;
    pid_t pid;
    char * argv[MAX_LENGTH];
    int status;

```

[illegible]

```

        }
    }
    if ( background )
    {
        pthread_t backgroundWaitingThread;
        PROCESS_STRUCTURE processStructure;
        processStructure.pid = pid;
        strncpy( processStructure.cmdline, argv[0] , MAXLENGTH )
        pthread_create( &backgroundWaitingThread ,
                        NULL, waitBackgroundProcess ,
                        (void *)&processStructure );
        return true;
    }
    else
    {
        waitpid( pid , &status , 0 );
        if ( WEXITSTATUS(status) == EXIT_FAILURE )
            return false;
        else
            return true;
    }
}

int main (int argc , char * argv[])
{
    FILE * fp = NULL;
    bool debug = false;
    char cmd[MAXLENGTH];
    memset( cmd, 0, sizeof(cmd) );
    if ( argc >= 2 )
    {
        if ( !strcmp("-v", argv[1] ) )
            debug = true;
        else if ( !strcmp("--help", argv[1] ) )
        {
            arghelp(argv[0]);
            return EXIT_SUCCESS;
        }
        else if ( !strcmp("--version", argv[1] ) )
        {
            version();
            return EXIT_SUCCESS;
        }
        else

```

```

        {
            if ( ( fp = fopen(argv[1] , "r" ) ) == NULL )
            {
                printf( "victorshell: %s: invalid option\n",
                        argv[0] );
                return EXIT_FAILURE;
            }
        }
    }
    if ( fp == NULL )
        fp = stdin;
    while (true)
    {
        char path[MAXLENGTH];
        char hostname[MAXLENGTH];
        if ( fp == stdin )
        {
            gethostname( hostname , sizeof(hostname) );
            getcwd( path , sizeof(path) );
            printf ( "[%s@%s %s]" , getpwuid( getuid() )->pw_name,
                        hostname, path );
            if ( getuid() == 0 )
                printf ( "# " );
            else
                printf ( "$ " );
            fflush( stdout );
        }
        if ( fgets(cmd, sizeof(cmd)-1, fp ) == NULL )
        {
            if ( fp != stdin )
                fclose( fp );
            return EXIT_SUCCESS;
        }
        cmd[ strlen(cmd)-1 ] = '\0';
        if (debug)
            puts(cmd);
        if ( !exeInternalCmd(cmd) )
        {
            if( !exeExternalCmd(cmd) )
            {
                char argfirst[MAXLENGTH];
                sscanf( cmd, "%s", argfirst );
                printf( "victorshell: %s: command not found\n",
                        argv[0] );
            }
        }
    }
}

```

```
    }  
    return EXIT_SUCCESS;  
}
```

C Source Code of Synchronization

C.1 process_synchronization.c

```
/*
 * process_synchronization.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *     gcc process_synchronization.c -o \
 *         process_synchronization -std=gnu99 -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <time.h>
#include <stdbool.h>
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/sem.h>
#include <unistd.h>

typedef int COUNT;

union semun {
    int          val;          /* Value for SETVAL */
    struct semid_ds * buf;      /* Buffer for IPC_STAT, IPC_SET */
    unsigned short * array;     /* Array for GETALL, SETALL */
    struct seminfo * __buf;     /* Buffer for IPC_INFO
                                (Linux-specific) */
};

int createSemaphore ( void )
{
    int semid;
    union semun arg;
    semid = semget( IPC_PRIVATE, 1, 0666|IPC_CREAT );
    if ( semid == -1 )
        return -1;
    arg.val = 1;
    if ( semctl( semid, 0, SETVAL, arg ) == -1 )
```



```

        return -1;
    return semid;
}

bool destroySemaphore ( int semid )
{
    union semun arg;
    if ( semctl( semid, 0, IPC_RMID, arg ) == -1 )
        return false;
    else
        return true;
}

bool P( int semid )
{
    struct sembuf spos;
    spos.sem_num = 0;
    spos.sem_op = -1;
    spos.sem_flg = 0;
    if ( semop( semid, &spos, 1 ) == -1 )
        return false;
    return true;
}

bool V( int semid )
{
    struct sembuf spos;
    spos.sem_num = 0;
    spos.sem_op = 1;
    spos.sem_flg = 0;
    if ( semop( semid, &spos, 1 ) == -1 )
        return false;
    return true;
}

int main (void)
{
    const int n = 10;
    int maxLoop[7];
    COUNT i;
    pid_t pid;

    int p12 = createSemaphore();
    int p13 = createSemaphore();

```

```

int p24 = createSemaphore();
int p25 = createSemaphore();
int p35 = createSemaphore();
int p46 = createSemaphore();
int p56 = createSemaphore();
int p6  = createSemaphore();

srand( time(NULL) );
for ( i = 1; i < 7 ; i ++ )
    maxLoop[i] = (rand() % n) + 1;
P(p12);
P(p13);
P(p24);
P(p25);
P(p35);
P(p46);
P(p56);
P(p6);

for ( i = 1; i <= 6 ; i ++ )
{
    if ( (pid = fork() ) < 0 )
    {
        fprintf( stderr , "fork()_error!\n\n" );
        return EXIT_FAILURE;
    }
    else if ( pid == 0 )
        break;
}

switch ( i )
{
    case 1:
        for ( COUNT j = 0 ; j < maxLoop[1] ; j ++ )
        {
            puts( "I_am_process_one." );
            usleep( 500000 );
        }
        V(p12);
        V(p13);
        return EXIT_SUCCESS;
    case 2:
        P(p12);
        for ( COUNT j = 0 ; j < maxLoop[2] ; j ++ )

```

```

        {
            puts( "I_am_process_two." );
            usleep( 500000 );
        }
        V(p24);
        V(p25);
        return EXIT_SUCCESS;
case 3:
    P(p13);
    for ( COUNT j = 0 ; j < maxLoop[3] ; j ++ )
    {
        puts( "I_am_process_three." );
        usleep( 500000 );
    }
    V(p35);
    return EXIT_SUCCESS;
case 4:
    P(p24);
    for ( COUNT j = 0 ; j < maxLoop[4] ; j ++ )
    {
        puts( "I_am_process_four." );
        usleep( 500000 );
    }
    V(p46);
    return EXIT_SUCCESS;
case 5:
    P(p25);
    P(p35);
    for ( COUNT j = 0 ; j < maxLoop[5] ; j ++ )
    {
        puts( "I_am_process_five." );
        usleep( 500000 );
    }
    V(p56);
    return EXIT_SUCCESS;
case 6:
    P(p46);
    P(p56);
    for ( COUNT j = 0 ; j < maxLoop[6] ; j ++ )
    {
        puts( "I_am_process_six." );
        usleep( 500000 );
    }
    V(p6);

```

```

        return EXIT_SUCCESS;
default:
    P(p6);
    destroySemaphore( p12 );
    destroySemaphore( p13 );
    destroySemaphore( p24 );
    destroySemaphore( p25 );
    destroySemaphore( p35 );
    destroySemaphore( p46 );
    destroySemaphore( p56 );
    destroySemaphore( p6 );
    puts("All Processes have finished.");
    return EXIT_SUCCESS;
}

return EXIT_SUCCESS;
}

```

C.2 thread_synchronization.c

```
/*
 * thread_synchronization.c
 *
 * Copyright Victor Jianfei Ye 2012
 *
 * Compile command:
 *     gcc thread_synchronization.c -o \
 *         thread_synchronization -std=gnu99 \
 *         -D_REENTRANT -lpthread -Wall
 */

#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <pthread.h>
#include <unistd.h>
#include <semaphore.h>
#include <time.h>

typedef int COUNT;

sem_t p12, p13, p24, p25, p35, p46, p56;
int maxLoop[7];

void * thread1 ( void * args )
{
    for ( COUNT i = 0 ; i < maxLoop[1] ; i ++ )
    {
        puts( "I am thread one." );
        usleep(500000);
    }
    sem_post( &p12 );
    sem_post( &p13 );
    return NULL;
}

void * thread2 ( void * args )
{
    sem_wait( &p12 );
    for ( COUNT i = 0 ; i < maxLoop[2] ; i ++ )
```

```

        {
            puts( "I_am_thread_two." );
            usleep(500000);
        }
        sem_post( &p24 );
        sem_post( &p25 );
        return NULL;
    }

void * thread3 ( void * args )
{
    sem_wait( &p13 );
    for ( COUNT i = 0 ; i < maxLoop[3] ; i ++ )
    {
        puts( "I_am_thread_three." );
        usleep(500000);
    }
    sem_post( &p35 );
    return NULL;
}

void * thread4 ( void * args )
{
    sem_wait( &p24 );
    for ( COUNT i = 0 ; i < maxLoop[4] ; i ++ )
    {
        puts( "I_am_thread_four." );
        usleep(500000);
    }
    sem_post( &p46 );
    return NULL;
}

void * thread5 ( void * args )
{
    sem_wait( &p25 );
    sem_wait( &p35 );
    for ( COUNT i = 0 ; i < maxLoop[5] ; i ++ )
    {
        puts( "I_am_thread_five." );
        usleep(500000);
    }
    sem_post( &p56 );
    return NULL;
}

```

```

}

void * thread6 ( void * args )
{
    sem_wait( &p46 );
    sem_wait( &p56 );
    for ( COUNT i = 0 ; i < maxLoop[6] ; i ++ )
    {
        puts( "I_am_thread_six." );
        usleep(500000);
    }
    return NULL;
}

int main (void)
{
    pthread_t threadid[7];
    const int n = 10;

    srand( time(NULL) );

    for ( COUNT i = 1 ; i < 7 ; i ++ )
        maxLoop[i] = (rand() % n) + 1;

    if ( sem_init( &p12, 0, 0 ) == -1 )
    {
        fprintf( stderr , "Error_Initializing_Semaphore\n\n");
        return EXIT_FAILURE;
    }
    if ( sem_init( &p13, 0, 0 ) == -1 )
    {
        fprintf( stderr , "Error_Initializing_Semaphore\n\n");
        return EXIT_FAILURE;
    }
    if ( sem_init( &p24, 0, 0 ) == -1 )
    {
        fprintf( stderr , "Error_Initializing_Semaphore\n\n");
        return EXIT_FAILURE;
    }
    if ( sem_init( &p25, 0, 0 ) == -1 )
    {
        fprintf( stderr , "Error_Initializing_Semaphore\n\n");
        return EXIT_FAILURE;
    }
}

```

```

if ( sem_init( &p35, 0, 0 ) == -1 )
{
    fprintf( stderr , "Error_Initializing_Semaphore\n\n");
    return EXIT_FAILURE;
}
if ( sem_init( &p46, 0, 0 ) == -1 )
{
    fprintf( stderr , "Error_Initializing_Semaphore\n\n");
    return EXIT_FAILURE;
}
if ( sem_init( &p56, 0, 0 ) == -1 )
{
    fprintf( stderr , "Error_Initializing_Semaphore\n\n");
    return EXIT_FAILURE;
}

if ( pthread_create( &threadid[1], NULL, thread1 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );
    return EXIT_FAILURE;
}
if ( pthread_create( &threadid[2], NULL, thread2 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );
    return EXIT_FAILURE;
}
if ( pthread_create( &threadid[3], NULL, thread3 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );
    return EXIT_FAILURE;
}
if ( pthread_create( &threadid[4], NULL, thread4 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );
    return EXIT_FAILURE;
}
if ( pthread_create( &threadid[5], NULL, thread5 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );
    return EXIT_FAILURE;
}
if ( pthread_create( &threadid[6], NULL, thread6 , NULL ) )
{
    fprintf( stderr , "Error_Creating_Thread.\n\n" );

```



```
        return EXIT_FAILURE;
    }

    for ( COUNT i = 1 ; i < 7 ; i ++ )
    {
        if ( pthread_join( threadid[i] , NULL ) )
        {
            fprintf( stderr , "Error_Joining_Thread.\n\n" );
            return EXIT_FAILURE;
        }
    }

    return EXIT_SUCCESS;
}
```
